

Introduction to R Programming

Slide Set 2: Exploratory Data Analysis with R

Maria Ptashkina

Barcelona GSE ITFD

September 2021

Table of Contents

① Before We Begin

② Data Wrangling

③ Importing, Exporting and Merging

- When you start working, you need to keep your files somewhere
- In R you need to declare a working directory
- A better way (which we will use for the applications next week) is using R Projects [▶ Projects](#)
- If you decide to use R for your Master Projects, you'd need to collaborate on the code. For this purpose I recommend doing it with Dropbox and creating R projects [▶ Dropbox and R](#)

Workflow

The screenshot displays the RStudio environment. The 'Session' menu is open, showing options like 'Interrupt R...', 'Restart R', 'Set Working Directory', 'Load Workspace...', 'Save Workspace As...', 'Clear Workspace...', and 'Quit Session...'. The 'Set Working Directory' option is highlighted, with a submenu showing 'To Source File Location', 'To Files Pane Location', and 'Choose Directory...'. The console shows the following R code:

```
##  
##  
##  
#####  
##  
#  
getwd()  
##  
setwd("C:/Users/Usua  
##  
# see the error  
qplot  
##  
# install the package  
install.packages("ggplot2")  
##  
# load the package / library  
library(ggplot2)
```

The console output shows the current working directory is 'C:/Users/Usua/Desktop/r course/_codes/' and the package 'ggplot2' is loaded. The file explorer on the right shows the 'Home' directory with various files and folders.

Packages

- So far we used 'base packages' which are part of R source code
- Most of the commands are part of the packages created by R users (there are 10,000+ user contributed packages and growing)
- Packages are collections of R functions, data, and compiled code in a well-defined format, created to add specific functionality
- You need to install a package only once using `install.packages("nameofpackage")` at the command line
- Then you need to call it in script using `library(nameofpackage)` every time you need to use it in a new R session
- If you'd like to use a specific command from a specific package use `package::command`

Getting Help

- Each R function comes with its own help page
- Type question mark and name of the function `?sqrt`
- If you forgot the function's name, you can search my keyword using two question marks `??log`
- This is all very formal and nice, but in reality you will literally only use Google and Stack Overflow

Functions

- Before we dive into data analysis, let's refresh a few key concepts
- Remember that R works with objects that you create
- The job is to apply functions to objects
- A standard way to work in R is to nest functions

`round(mean(die))`



`round(mean(1:6))`




`round(3.5)`



4

Tidyverse

- Tidyverse is a collection of integrated packages which was designed to make working with data more user friendly
- You can do all the same with nested functions and individual packages, if you prefer
- Tidyverse helps you to clean the data and put it into a format which you would need to build a model / run a regression



Pipes

- Stringing together commands can be difficult, especially if there are many nested functions
- The pipe allows the output of a previous command to be used as input to another command instead of using nested functions
- Shortcut: Shift+Ctrl+M

%>%

```
leave_house(get_dressed(get_out_of_bed(wake_up(me, time =  
"8:00"), side = "correct"), pants = TRUE, shirt = TRUE), car  
= TRUE, bike = FALSE)
```

```
me %>%  
  wake_up(time = "8:00") %>%  
  get_out_of_bed(side = "correct") %>%  
  get_dressed(pants = TRUE, shirt = TRUE) %>%  
  leave_house(car = TRUE, bike = FALSE)
```

- A core component of the tidyverse is the tibble
- It's basically a modern (more flexible and efficient) version of a data frame
- Tibbles can be created directly using the `tibble()`
- Data frames can be converted into tibbles using `as_tibble(df)`


Table of Contents

① Before We Begin

② Data Wrangling

③ Importing, Exporting and Merging

Plan



Data types in Tibbles

- `int` stands for integers
- `dbl` stands for doubles, or real numbers
- `chr` stands for character vectors, or strings
- `dtm` stands for date-times (a date + a time)
- `lgl` stands for logical, vectors that contain only TRUE or FALSE
- `fctr` stands for factors, which R uses to represent categorical variables with fixed possible values
- `date` stands for date

- The vast majority of your data manipulation challenges are going to be handled using `dplyr` package (part of tidyverse):
 - Pick observations by their values `filter()`
 - Reorder the rows `arrange()`
 - Pick variables by their names `select()`
 - Create new variables with functions of existing variables `mutate()`
 - Collapse many values down to a single summary `summarise()`
- These can all be used in conjunction with `group_by()` which changes the scope of each function

- `filter`: subset observations based on their values
- Use logical operators we saw before
- Missing values: represented in R as NA
 - Missing values are “contagious”: almost any operation involving an unknown value will also be unknown
 - Check for NA using `is.na()`
 - Note! R treats missing values differently from other statistical programs
 - ▶ [More Information](#)
- As we discussed, the majority of tidyverse has equivalents in base R
 - `filter()` is equivalent to `subset()`
 - ▶ [Difference](#)

Arrange

- `arrange()`: changes the order or rows
- Missing values are always sorted at the end
- As before you can use an equivalent base R command `order` (note how much less intuitive the syntax is) [▶ Order](#)

- `select()`: select the columns you need
- Helper functions
 - `starts_with("abc")`: matches names that begin with "abc"
 - `ends_with("xyz")`: matches names that end with "xyz"
 - `contains("ijk")`: matches names that contain "ijk"
 - `num_range("x", 1:3)`: matches x1, x2 and x3
- `rename()` is a variant of `select()`
- `select()` with `everything()` helper to move columns to the beginning
- Remember we were sub-setting vectors and matrices and data frames? You can do exactly the same as select with base R [▶ Sub-setting](#)

- `mutate()`: add new columns that are functions of existing columns
- If you only want to keep the new variables, use `transmute()`
- The function must be vectorised
- Frequently used functions
 - Arithmetic operators `+`, `-`, `*`, `/`, `^`
 - Logarithmic transformation
 - Offsets `lead()` and `lag()`
 - Cumulative and rolling aggregates
 - Logical comparisons, `<`, `<=`, `>`, `>=`, `!=`, `==`

- `summarise()`: collapse a data frame to a single row generating a statistic
- Usually combined with `group_by()`
- Frequently used functions
 - Measures of location: mean, median, etc.
 - Measures of spread: standard deviation, median absolute deviation, etc.
 - Measures of rank: min, max, etc.
 - Counts: size of a particular group, number of unique values
 - Counts and proportions of logical values (when used with numeric functions, TRUE is converted to 1 and FALSE to 0)

Table of Contents

① Before We Begin

② Data Wrangling

③ Importing, Exporting and Merging

Importing Data

Data type	Extension	Function	Package
Coma separated values	csv	read.csv()	base
		read_csv()	readr (tidyverse)
Other delimited formats	txt	read.table()	base
		read_table()	readr (tidyverse)
		read_delim()	readr (tidyverse)
Excel	xlsx, xls	read_excel()	readxl (tidyverse)
Stata version 13 and above	dta	read_dta()	haven
Stata version 7-12	dta	read.dta()	foreign
SPSS	sav	read.spss()	foreign

It's advised to use tidyverse equivalents, because the functions


- Are x10 times faster
- Don't convert character vectors to factors
- Are more reproducible

Exporting Data


- Similarly, you can export the data to other formats, using `write.table`, `write.csv` or `write_csv`, and `write.dta`
- You can also save data into R data format: RDATA and RDS
- You can read about other ways of exporting data [▶ Exporting](#)

- Keys: the variables used to connect each pair of tables and *uniquely* identifies an observation
 - Primary key: uniquely identifies an observation in its own table
 - Foreign key uniquely identifies an observation in another table
- Relations are typically one-to-many or many-to-one, but sometimes are one-to-one
- Types of joins
 - Inner join
 - Outer joins: left, right, full


Inner Join




Left Join



Right Join



Full Join



Defining Key Columns

- If your data has many levels of variation, you might need to join by several variables
- By default, `by = NULL` uses all variables that appear in both tables
- You can define only some variables using a character vector
- You can also match variable *a* in table *x* to variable *b* in table *y*

Comparing with Base R

dplyr	base
<code>inner_join(x, y)</code>	<code>merge(x, y)</code>
<code>left_join(x, y)</code>	<code>merge(x, y, all.x = TRUE)</code>
<code>right_join(x, y)</code>	<code>merge(x, y, all.y = TRUE)</code>
<code>full_join(x, y)</code>	<code>merge(x, y, all.x = TRUE, all.y = TRUE)</code>

- In base R you can also merge by character vectors using `by.x=c()` and `by.y=()`

- A handy tool for data inspection is to join by observations as opposed to variables
 - `semi_join(x, y)` keeps all observations in `x` that have a match in `y`
 - `anti_join(x, y)` drops all observations in `x` that have a match in `y`
- Anti-joins are useful for diagnosing join mismatches
- Be aware that in your research the data might be much nastier, so explore it carefully before!

References and Resources

- Introduction to R Flipped ▶ Tutorial
- Hands-On Programming with R ▶ Tutorial
- R for Data Science ▶ Tutorial
- Tidyverse in R – Complete Tutorial ▶ Tutorial
- Data Wrangling with Tidyverse ▶ Tutorial
- Pipes ▶ Pipes
- Parsing a File (chapter 11.4) ▶ Parsing
- Getting Started with the Tidyverse: Tutorial ▶ Parsing